


Chapter 8: Main Memory





Chapter 8: Memory Management

- Background
- Swapping
- Contiguous Memory Allocation
- Segmentation
- Paging
- Structure of the Page Table





Objectives

- To provide a detailed description of various ways of organizing memory hardware
- To discuss various memory-management techniques, including paging and segmentation
- To provide a detailed description of the Intel Pentium, which supports both pure segmentation and segmentation with paging





Background

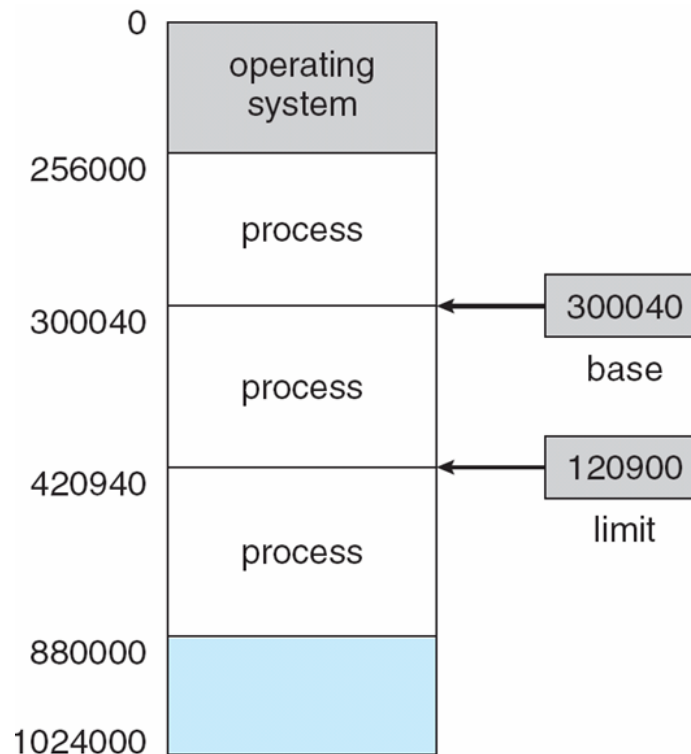
- Main memory and registers are the only storage CPU can access directly
- Program must be brought (from disk) into memory for it to be run
- Memory unit only sees a stream of addresses + read requests, or address + data and write requests
- Register access in one CPU clock
- Main memory can take many cycles, causing a **stall**
- **Cache** sits between main memory and CPU registers
- Protection of memory required to ensure correct operation
 - Protect OS from user programs
 - Protect user programs from each other





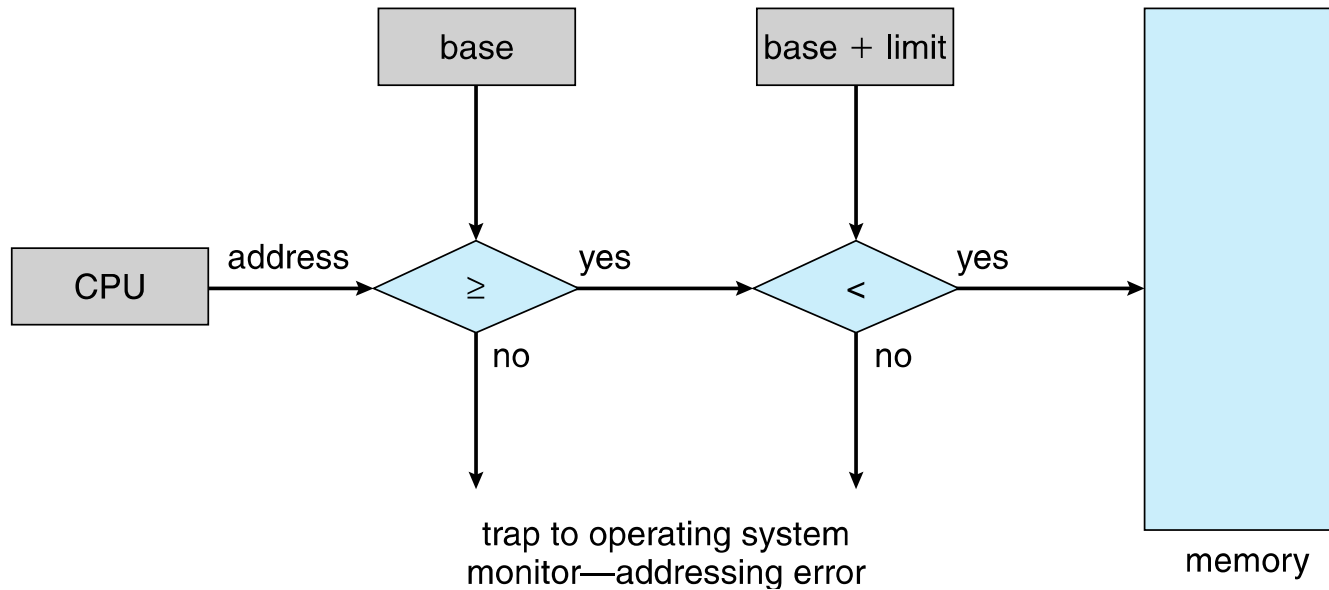
Base and Limit Registers

- A pair of **base** and **limit registers** define the logical address space
- CPU must check every memory access generated in user mode to be sure it is between base and limit for that user
- Only OS can load base and limit regs, using privileged instructions





Hardware Address Protection





Address Binding

- Programs on disk, ready to be brought into memory to execute form an **input queue**
- Inconvenient to have first user process physical address always at 0000
- Further, addresses represented in different ways at different stages of a program's life
 - Source code addresses usually symbolic
 - Compiled code addresses **bind** to relocatable addresses
 - ▶ i.e. "14 bytes from beginning of this module"
 - Loader will bind relocatable addresses to absolute addresses
 - ▶ i.e. 74014
 - Each binding maps one address space to another





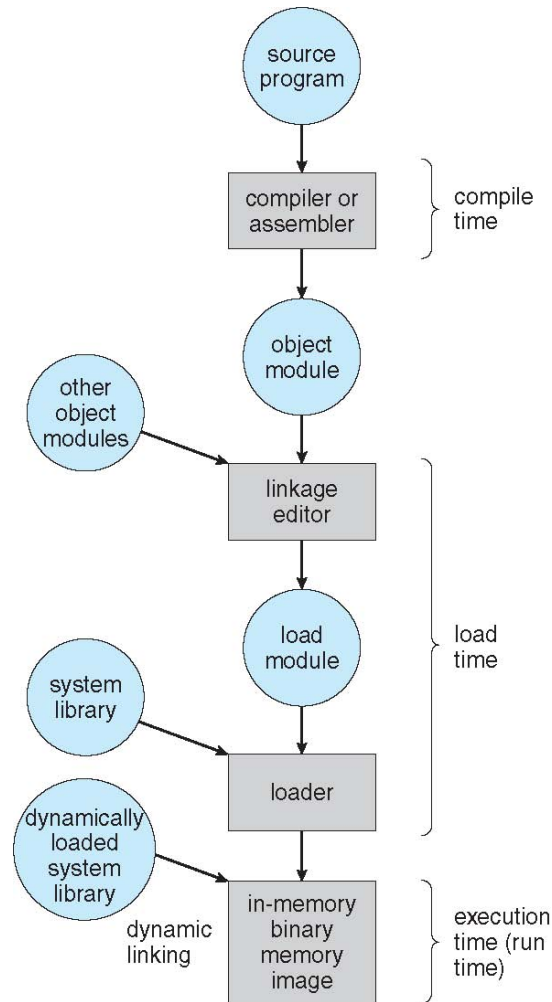
Binding of Instructions and Data to Memory

- Address binding of instructions and data to memory addresses can happen at three different stages
 - **Compile time:** If memory location known a priori, **absolute code** can be generated; must recompile code if starting location changes
 - **Load time:** Must generate **relocatable code** if memory location is not known at compile time
 - **Execution time:** Binding delayed until run time if the process can be moved during its execution from one memory segment to another
 - ▶ Need hardware support for address maps (e.g., base and limit registers)





Multistep Processing of a User Program





Logical vs. Physical Address Space

- The concept of a logical address space that is bound to a separate **physical address space** is central to proper memory management
 - **Logical address** – generated by the CPU; also referred to as **virtual address**
 - **Physical address** – address seen by the memory unit
- Logical and physical addresses are the same in compile-time and load-time address-binding schemes
- Logical (virtual) and physical addresses differ in execution-time address-binding scheme
- **Logical address space** is the set of all logical addresses generated by a program
- **Physical address space** is the set of all physical addresses generated by a program






Memory-Management Unit (MMU)

- Hardware device that at run time maps a virtual address to a physical address
- To start, consider simple scheme where the value in the relocation register is added to every address generated by a user process at the time it is sent to memory
 - Base register now called **relocation register**
- The user program deals with *logical* addresses; it never sees the *real* physical addresses
 - Execution-time binding occurs when reference is made to location in memory
 - Logical address bound to physical addresses





Dynamic linking using a relocation register





Dynamic Linking

- **Static linking:** system libraries and program code combined by the linker and the loader into the binary program image
- **Dynamic linking:** linking postponed until execution time
- System libraries also known as **shared libraries**
- Small piece of code, **stub**, used to locate the appropriate memory-resident library routine or load it if not already present
- Stub replaces itself with the address of the routine, and executes the routine
- Next time the routine is referenced, no dynamic linking cost
- Dynamic linking is particularly useful for libraries
- Consider applicability to patching system libraries
 - Versioning may be needed





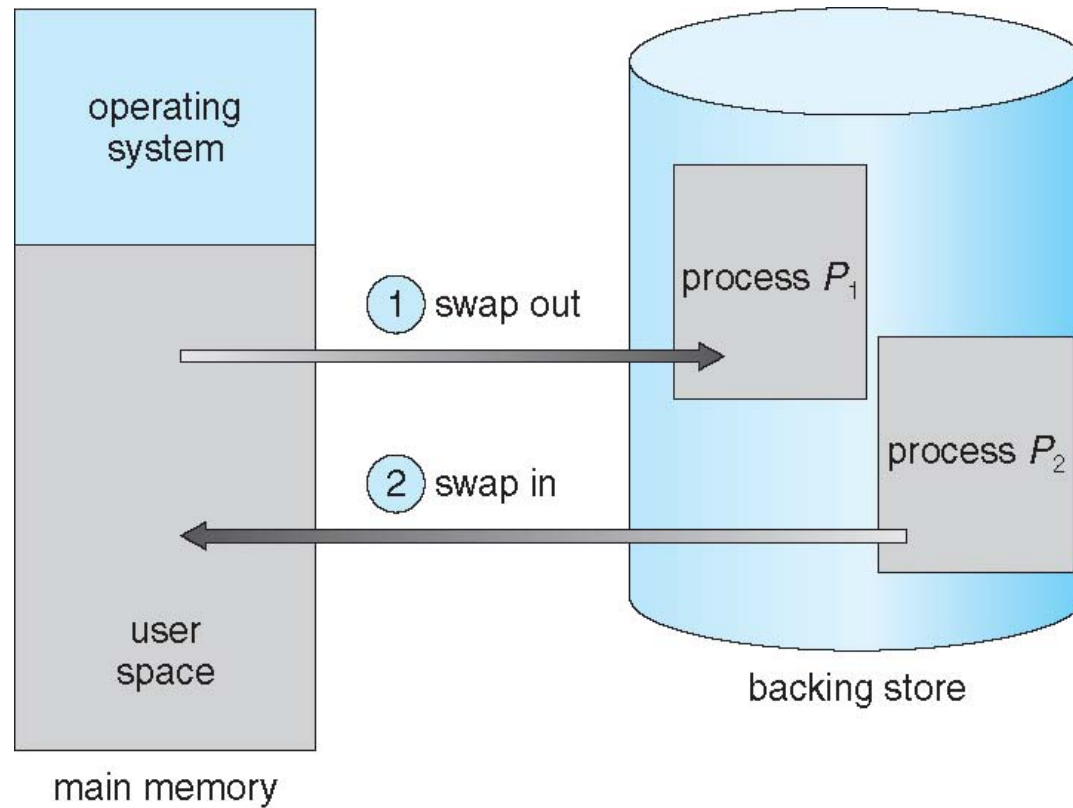
Swapping

- Total physical memory space of processes can exceed physical memory
- A process can be **swapped** temporarily out of memory to a backing store, and then brought back into memory for continued execution
- System maintains a **ready queue** of ready-to-run processes which have memory images on disk
- **Backing store** – fast disk large enough to accommodate copies of all memory images for all users; must provide direct access to these memory images
- Major part of swap time is transfer time; total transfer time is directly proportional to the amount of memory swapped





Schematic View of Swapping





Problems with Swapping

- If next processes to be put on CPU is not in memory, need to swap out a process and swap in target process
- Context switch time can then be very high
- 100MB process swapping to hard disk with transfer rate of 50MB/sec
 - Swap out time of 2000 ms
 - Plus swap in of same sized process
 - Total context switch swapping component time of 4 seconds
- Other constraints as well on swapping
 - Pending I/O – can't swap out as I/O would occur to wrong process
 - Or always transfer I/O to kernel space, then to process
 - ▶ Known as **double buffering**, adds overhead





Swapping in Modern Systems

- Standard swapping not used in modern operating systems
- Modified versions of swapping are found on many systems (i.e., UNIX, Linux, and Windows)
 - Swapping normally disabled
 - Started if memory demand exceeds a certain threshold
 - Disabled again once memory demand drops below threshold





Contiguous Allocation

- Main memory must support both OS and user processes
- Limited resource, must allocate efficiently
- Contiguous allocation is one early method
- Main memory usually into two **partitions**:
 - Resident operating system, usually held in low memory with interrupt vector
 - User processes held in high memory
 - Each process contained in single contiguous section of memory





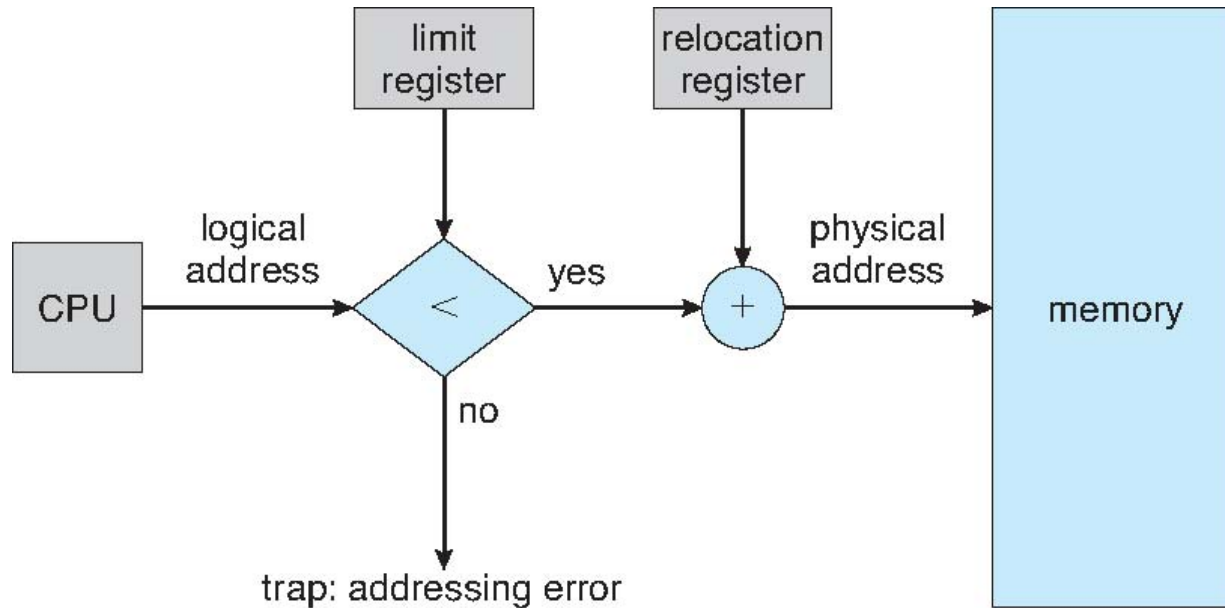
Contiguous Allocation (Cont.)

- Relocation registers used to protect user processes from each other and from changing operating-system code and data
 - Base register contains value of smallest physical address
 - Limit register contains range of logical addresses – each logical address must be less than the limit register
 - MMU maps logical address *dynamically*
 - Correct values are loaded into relocation and limit registers as part of context switching
 - Can then allow actions such as changing kernel size dynamically
 - ▶ E.g. Code that is not commonly used can be made **transient**





Hardware Support for Relocation and Limit Registers

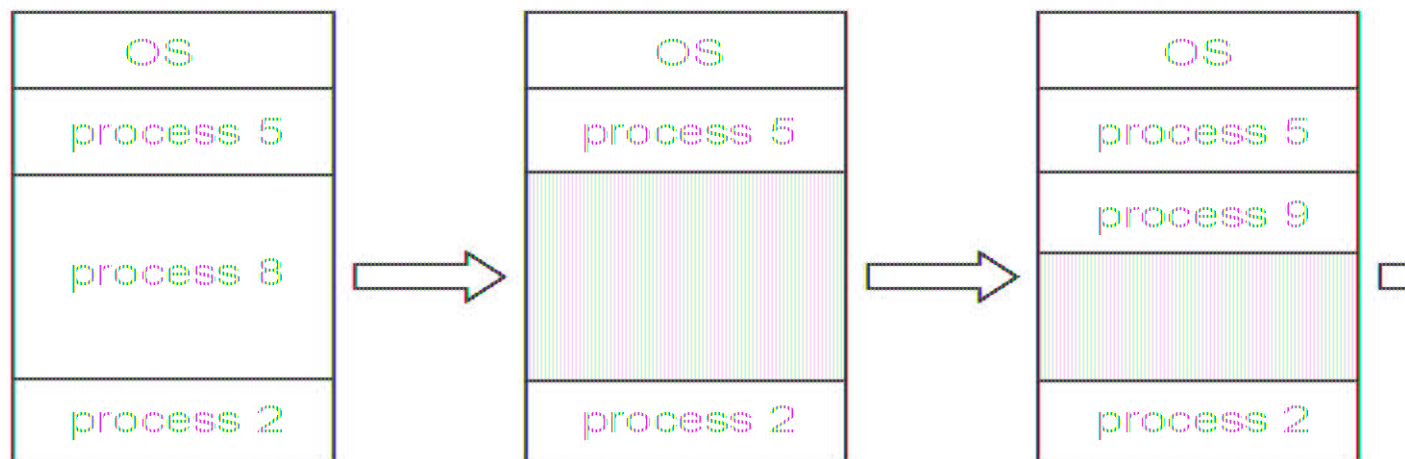




Multiple-partition Allocation

■ Multiple-partition allocation

- Degree of multiprogramming limited by number of partitions
- **Variable-partition** for efficiency (sized to a given process' needs)
- **Hole** – block of available memory; holes of various size are scattered throughout memory
- When a process arrives, it is allocated memory from a hole large enough to accommodate it
- Process exiting frees its partition, adjacent free partitions combined
- Operating system maintains information about:
a) allocated partitions b) free partitions (holes)





Dynamic Storage-Allocation Problem

How to satisfy a request of size n from a list of free holes?

- **First fit:** Allocate the ***first*** hole that is big enough
- **Best fit:** Allocate the ***smallest*** hole that is big enough; must search entire list, unless ordered by size
 - Produces the smallest leftover hole
- **Worst fit:** Allocate the ***largest*** hole; must also search entire list
 - Produces the largest leftover hole

First fit and best fit better than worst-fit in terms of storage utilization

First-fit is faster than best-fit





Fragmentation

- **External Fragmentation** – total memory space exists to satisfy a request, but it is not contiguous
- First fit analysis reveals that given N blocks allocated, $0.5 N$ blocks lost to fragmentation
 - $1/3$ may be unusable -> **50-percent rule**
- **Internal Fragmentation** – allocated memory may be slightly larger than requested memory; this size difference is memory internal to a partition, but not being used





Fragmentation (Cont.)

- Reduce external fragmentation by **compaction**
 - Shuffle memory contents to place all free memory together in one large block
 - Compaction is possible *only* if relocation is dynamic, and is done at execution time
- More effective solutions to fragmentation permit logical address space of a process to be non-contiguous
 - Segmentation
 - Paging






Segmentation

- Memory-management scheme that supports user view of memory
- A program is a collection of segments
 - A compiler might create separate segments for the following:
 - The code
 - Global Variables
 - The heap (dynamic memory allocation)
 - Stacks used by threads
 - Language Library






User's View of a Program






Logical View of Segmentation



user space

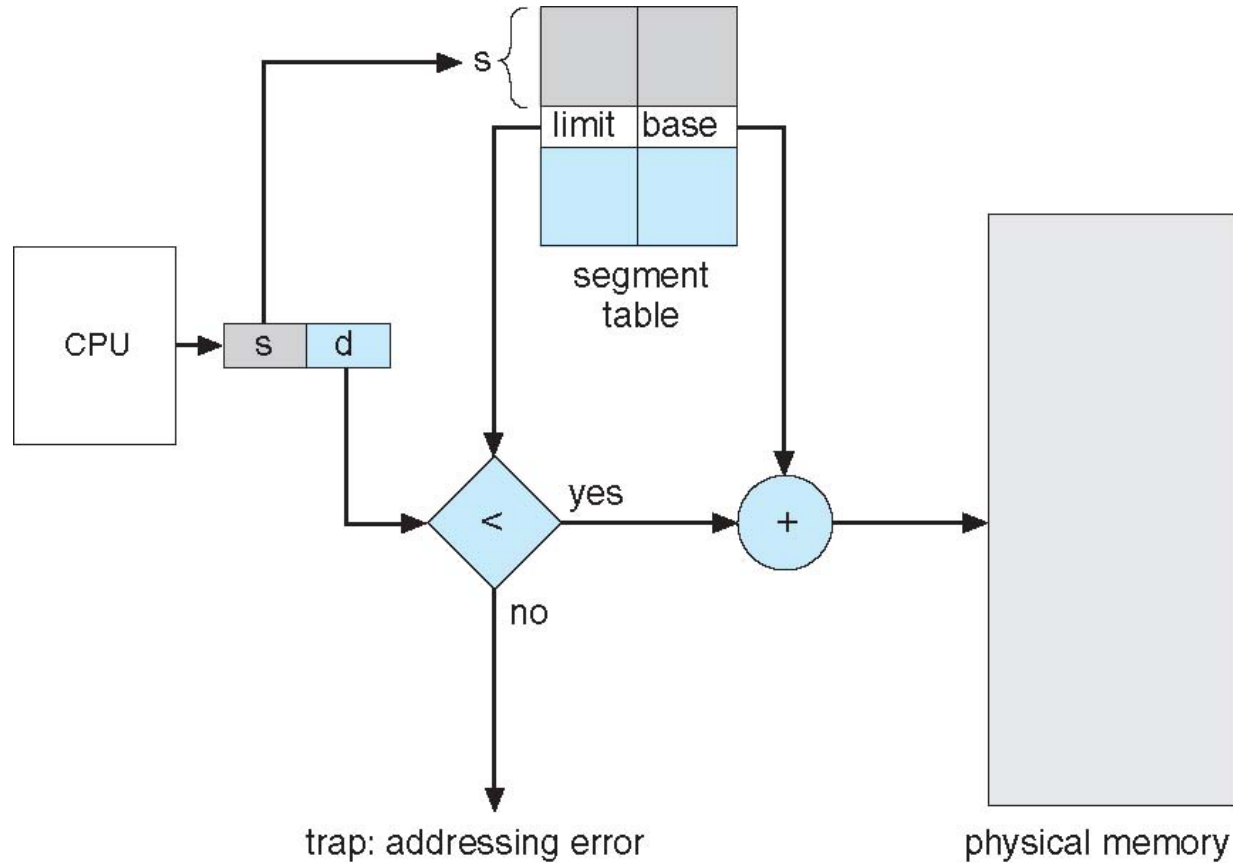


physical memory space





Segmentation Hardware





Segmentation Architecture

- Logical address consists of a two tuple:
 <segment-number, offset>,
- **Segment table** – maps two-dimensional user addresses into one-dimensional physical addresses; each table entry has:
 - **base** – contains the starting physical address where the segment resides in memory
 - **limit** – specifies the length of the segment





Paging


- Physical address space of a process can be noncontiguous; process is allocated physical memory whenever the latter is available
 - Avoids external fragmentation and the need for compaction
 - Avoids problem of varying sized memory chunks
 - Still has internal fragmentation
- Divide physical memory into fixed-sized blocks called **frames**
 - Size is power of 2, between 512 bytes and 16 Mbytes
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of size ***N*** pages, need to find ***N*** free frames and load program
- Set up a **page table** to translate logical to physical addresses
- Backing store likewise split into pages





Address Translation Scheme

- Address generated by CPU is divided into:
 - **Page number** (p) – used as an index into a **page table** which contains base address of each page in physical memory
 - **Page offset** (d) – combined with base address to define the physical memory address that is sent to the memory unit

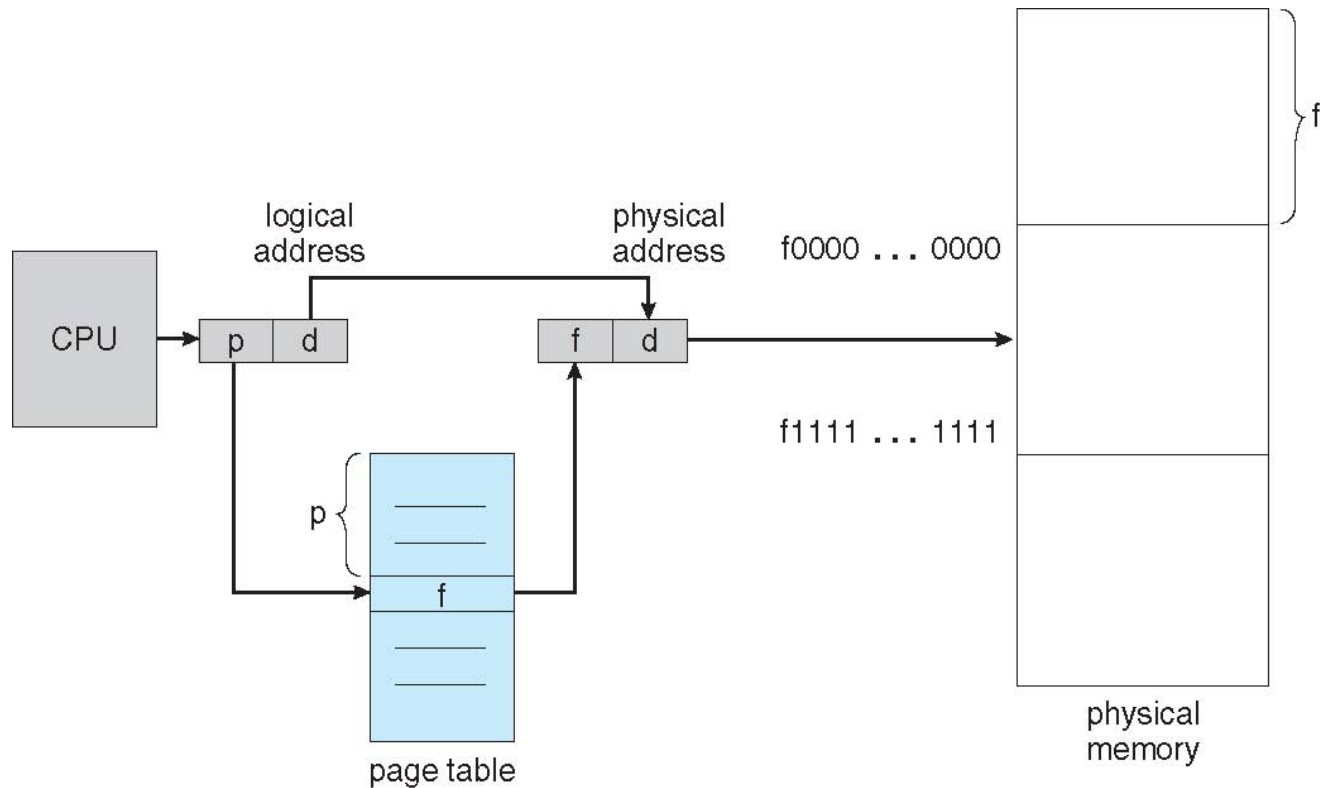


- For given logical address space 2^m and page size 2^n
 - ▶ E.g. $m=32$, $n=10$ (page size = 1K)






Paging Hardware





Paging Model of Logical and Physical Memory




logical
memory

0	1
1	4
2	3
3	7

page table

frame
number

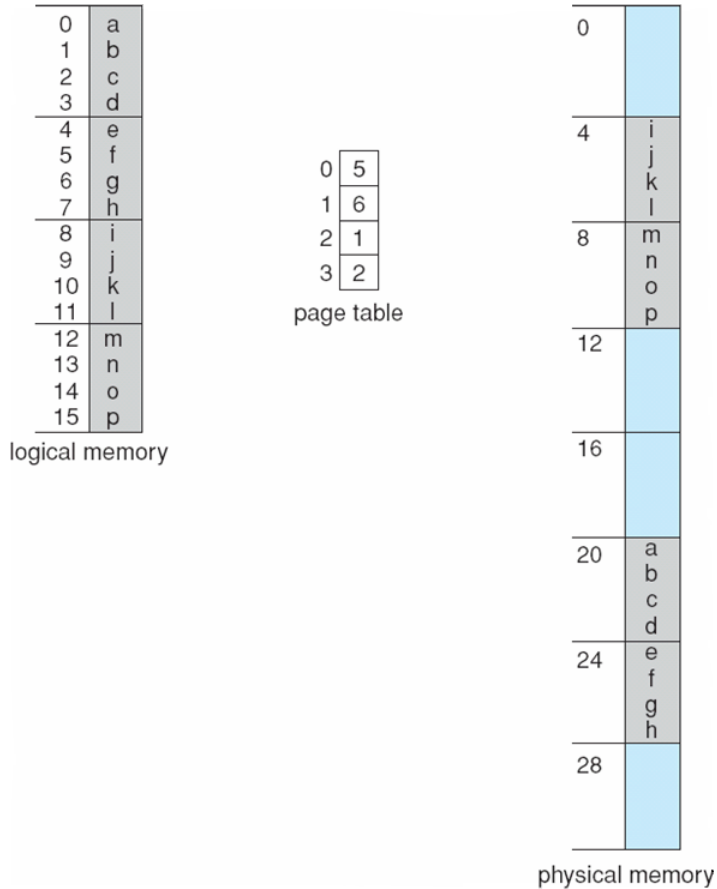


physical
memory





Paging Example



4-byte pages, 16-byte logical memory, 32-byte physical memory





Internal Fragmentation

- Calculating internal fragmentation
 - Page size = 512 bytes
 - Process size = 2060 bytes
 - 4 pages + 12 bytes
 - Internal fragmentation of $512 - 12 = 500$ bytes
 - Worst case fragmentation = 1 frame – 1 byte
 - On average fragmentation = $1 / 2$ frame size
 - So, using a smaller frame size reduces fragmentation, but each page table entry takes memory to track -> tradeoff!
 - Page sizes growing over time
 - ▶ Typically 4 KB to 8 KB
 - ▶ Some systems support larger frames






Addressing

- How much memory can we address with 4-byte page table entries?
 - 32-bit entry can address up to 2^{32} page frames
 - If the frame size is 4KB (2^{12}), then we can address 2^{44} bytes (16 TB)
- Actual size will be smaller, because other information will have to be kept in a page table entry
- Process view and physical memory now very different
 - Process view is contiguous, while it is actually scattered throughout physical memory
- By implementation, process can only access its own memory
 - A process has no way of addressing memory outside of its page table
 - Its page table includes only the frames that the process owns






Free Frames



(a)

Before allocation



(b)

After allocation





Implementation of Page Table

- If the page table is small enough, it can be implemented as registers
 - But, this is not feasible on most modern systems that need large page tables
- So, page table is kept in main memory
- **Page-table base register (PTBR)** points to the page table
- **Page-table length register (PTLR)** indicates size of the page table (number of entries)
- In this scheme every data/instruction access requires two memory accesses
 - One for the page table and one for the data / instruction
- The two-memory-access problem can be solved by the use of a special fast-lookup hardware cache called **translation look-aside buffer (TLB)** which is implemented using **associative memory (high speed)**.





Implementation of Page Table (Cont.)

- TLB has a key or tag (page number) and a value (frame number).
- TLBs typically small (32 to 1,024 entries)
- On a TLB miss, value is loaded into the TLB for faster access next time
 - Replacement policies must be considered
 - Some entries can be **wired down** for permanent fast access
- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
 - Otherwise need to flush at every context switch





Associative Memory

■ Associative memory – parallel search

P a g e #	F r a m e #


■ Address translation (p, d)

- If p is in associative register, get frame # out
- Otherwise get frame # from page table in memory





Paging Hardware With TLB





Effective Access Time

- Associative Lookup = ε time unit
 - Much smaller than memory access time
- Hit ratio = α
 - Hit ratio – percentage of times that a page number is found in the associative registers; ratio related to number of associative registers

- **Effective Access Time (EAT)**

$$\text{EAT} = (M + \varepsilon) \alpha + (2M + \varepsilon)(1 - \alpha)$$

M: memory access time

- Consider $\alpha = 80\%$, $\varepsilon = 2\text{ns}$ for TLB search (negligible), 100ns for memory access:
 - $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
- Consider more realistic hit ratio $\rightarrow \alpha = 99\%$, $\varepsilon = 2\text{ns}$ for TLB search, 100ns for memory access
 - $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns}$



End of Chapter 8

