

**Operating Systems Principles**  
**First Quiz, Section 1**  
**Friday, September 22<sup>nd</sup>, 2017**  
**Instructor: Ghassan Shobaki**

**Student Name:** \_\_\_\_\_ **Student Number:** \_\_\_\_\_

**Q1.** Modern operating systems are written in high-level languages (mainly C and C++). Give two good reasons for writing an operating system in a high-level language and explain each reason in the space provided. (14 points)

**Q2.** In a time sharing OS, one process runs into an infinite loop. Answer the following. (12 points)  
Does the OS detect infinite loops?

Does the OS terminate a process that runs into an infinite loop?

Will this cause other processes to hang? Why?

**Q3.** Compare the **monolithic** and the **loadable module** OS structures by giving a **clear** advantage of each structure relative to the other. Explain each advantage in the provided space. (14 points)  
Advantage of monolithic relative to loadable modules

Advantage of loadable modules relative to monolithic

**Q4.** The table below shows three events that may cause a transition in a process's state. For each event, the table shows the state transition that immediately follows the event and the queue in which the process gets placed after the transition. Fill in the missing entries in the table. If there is no queue, write **none**. (30 points)

| Causing Event                     | Transition |          | Queue after transition<br>(if none, enter <b>none</b> ) |
|-----------------------------------|------------|----------|---|
|                                   | From State | To State |   |
| System call to read from keyboard |            |          |   |
|                                   | Running    | Ready    |   |
| Scheduler dispatch                |            |          |   |

**Q5.** An OS running on a **single CPU** has the following two processes:

- $P_1$  consists of a 70-ms CPU burst followed by an I/O request followed by a 30-ms CPU burst
- $P_2$  consists of a 40-ms CPU burst followed by an I/O request followed by a 60-ms CPU burst

Draw a complete timing diagram showing how a time-sharing operating system will handle these processes under the following assumptions:

- Both I/O requests go to the same device and each request takes 80 ms.
- The I/O device will start processing a queued request as soon as the current request is done (no delay).
- Whenever  $P_1$  and  $P_2$  are **both** ready, the scheduler will give the CPU to the process that did not get the last CPU time (fair alternation if **both** are ready). In the beginning, the CPU is given to  $P_1$ .
- The time quantum is 50 ms
- The kernel time (including initiation of I/O requests, scheduling and context switching) is 10 ms.
- The OS starts each I/O request right at the beginning of the kernel time.
- Ignore any delay or transition time that was ignored in class. (30 points)

Show both the CPU and the I/O device in your diagram and make sure that you align the CPU time with the I/O time.

